

Few-shot learning using pre-training and shots, enriched by pre-trained samples

Detlef Schmicker

Ludgeriplatz 31, 47057 Duisburg, Germany

Abstract

We use the EMNIST dataset of handwritten digits to test a simple approach for few-shot learning. A fully connected neural network is pre-trained with a subset of the 10 digits and used for few-shot learning with untrained digits. Two basic ideas are introduced: during few-shot learning the learning of the first layer is disabled, and for every shot a previously unknown digit is used together with four previously trained digits for the gradient descend, until a predefined threshold condition is fulfilled. This way we reach about 90% accuracy after 10 shots.

1. Introduction

Neural networks have shown stunning success [1], but training requires a big database or the ability to generate a high amount of training samples from rules [2]. Obviously it is not always possible to fulfill this, for example if you want to teach a computer manually. Therefore we are interested to learn from few examples. A number of quite different approaches were proposed and it is even difficult to classify our approach within the discussed ones [3]. Our approach could be classified as a refinement of existing parameters, as well as embedding learning, no meta learning and one could argue that we use external memory, as we use samples from pre-training during few-shot learning.

We combine a number of well known concepts. To avoid optimizing the neural network for a specific task we use only fully connected layers with weights and without bias. The activation function has an output between 0 and 1 to keep it similar to human neurons (figure 1). The function value is close to 0 for input 0, similar to human neurons, where a 0 (not firing) has no effect on the connected neuron, as a zero output is only multiplied by a weight. Some tests indicate, that this restriction is not very important, but as it similar to human neurons, we keep it that way. This choice might also be useful, if one wants to add local plasticity [4]. In this work we use “plasticity” only by reducing the learning rate of the first layer.

The samples are taken from the EMNIST dataset of handwritten digits [5]. The EMNIST dataset contains 280000 digits from more than 500 different writers, classified using 10 labels, representing the digits. The subset belonging to 8 of the 10 digits are used for pre-training. The motivation is, that humans have seen a lot of lines and shapes during there live before they try to read digits. Therefore they have a pre-trained brain. After pre-training the neural network learns the two remaining digits, which

it has never seen before, from few examples with a few shot procedure, we will describe in detail.

The pre-training is done with a standard gradient descend. The few-shot learning is also done this way, but only uses one new sample at a time, and stops learning depending on a stop-criteria. Without any additional measures, this approach fails. Two measures were necessary to succeed:

- the learning for the first layer has to be disabled, or at least slowed down, and

- with every shot some previously known samples have to be added.

This helps the network remembering the old labels, similar to human learning: if you do not use old knowledge, you forget it. This way we reach about 90% accuracy with 10 shot learning. The jupyter notebook with the calculations is available [6, 7].

2. The neural network

Each fully connected layer has i inputs x_i and j outputs y_j , which are related by $y_j = \sum_i w_{ij} x_i$. Each activation layer takes k inputs x_k and has k outputs y_k with the relation $y_k = \frac{1}{1 - e^{-(3(x_k - 1))^2}}$, which is a sigmoid function, scaled and shifted in x-direction (figure 1). The neural network has one input layer, two hidden layers and an output layer as in figure 2, but the size of the layers are larger in the following. The loss function $l(y_i) = \frac{1}{2} \sum_i (y_i - \hat{y}_i)^2$ is used, with \hat{y} being the target value from the training data. Every output of the network corresponds to one of the labels (digits), with this target being 1 and all others 0. The gray scale inputs are scaled to the range between 0 and 1. Usually the weights are random initialized with a mean value of 0 and a standard derivative of 0.1 and clipped during training to a value between -1 and 1, similar to human neurons not firing with infinite intensity. The minimum of the loss function with respect to the weights is

Email address: dschmicker@physik.de (Detlef Schmicker)




Figure 1: The activation function used. Similar to human neurons the output is close to zero, if there is no input from the previous layer.




Figure 2: This graph shows a fully connected neural network, with 5 inputs, two hidden layers of size 4 and 3 outputs. Throughout this work we use this neural network with larger layers.

calculated using a gradient descend method implemented in python [7].

3. The pre-training procedure

From the EMNIST dataset of handwritten digits we chose 8 digits for the pre-training. With them a neural network with 28x28 input pixels, two hidden layers of size 64 and 10 outputs is trained. From the 10 outputs only 8 are used for pre-training, the other two are reserved for few-shot learning of the remaining two digits (the digits 2 and 8 in this work). We train for 100000 batches with a batch size of 1000. The EMNIST dataset of handwritten digits contains 280000 digits, which are split into 240000 digits for training and 40000 digits for testing. From the 240000 digits only 192000 digits belong to the 8 digits we use for pre-training. The pre-training procedure results in an accuracy of about 98%, measured with 1000 images from the testing dataset. A prediction is considered correct, if the output belonging to the correct digit has the highest value of all outputs.

4. The few-shot procedure

Two modifications as compared to the pre-training are used during few-shot learning:

1. learning for the first layer is disabled *and*
2. for every shot one new sample of the previously unknown digits is combined with 4 samples, which were part of pre-training.

With this mini-batch a gradient descend is performed until the two following stop criteria for the new sample are fulfilled:

1. the value of the output, corresponding to the correct label, is bigger than 0.3, *and*
2. the value of the output, corresponding to the correct label, is more than 1.5 times the value of the second highest output.

The samples used during pre-training are not taken into account for the stop criteria. For every shot we perform a gradient descend with two new samples, belonging to both labels not used during pre-training.

5. Experimental results

The result of the few-shot learning, measured using the testing data set, is shown in table 1. The first accuracy column reports the accuracy, taking only pre-trained digits into account, the second reports how good the new digits are recognized. The “accuracy only new digits” column reports the accuracy, if one tries to distinguish the new digits. Therefore it is just verified that the neurons corresponding to the additional digits show the higher value. The last column measures how good the neural network

shot	accuracy old digits	accuracy new digits	accuracy only new digits	overall accuracy
1	0.895	0.492	0.776	0.815
2	0.952	0.574	0.754	0.879
3	0.950	0.521	0.667	0.870
4	0.960	0.535	0.730	0.878
5	0.856	0.621	0.677	0.818
6	0.961	0.708	0.839	0.905
7	0.958	0.666	0.806	0.900
8	0.949	0.747	0.844	0.911
9	0.954	0.811	0.918	0.923
10	0.926	0.836	0.919	0.908

Table 1: Experimental results for few-shot learning. The pre-training is done with 100000 batches of size 1000. Learning of the first layer is disabled during few-shot learning, and every shot is enriched by four samples, known from pre-training.

shot	accuracy old digits	accuracy new digits	accuracy only new digits	overall accuracy
1	0.485	0.597	0.706	0.496
2	0.296	0.669	0.716	0.357
3	0.288	0.310	0.67	0.289
4	0.335	0.677	0.727	0.399
5	0.387	0.683	0.811	0.443
6	0.382	0.720	0.806	0.438
7	0.400	0.759	0.893	0.474
8	0.352	0.845	0.900	0.442
9	0.487	0.502	0.820	0.479
10	0.386	0.617	0.947	0.433

Table 2: Experimental results as in 1 , but without fixed first layer.

recognizes all 10 digits. A accuracy of around 90% seems quite convincing for few-shot learning of handwritten digits from many different writers.

We check both modifications from section 4 experimentally: first we do not disable learning for the first layer (table 2) and second we do not enrich the few-shot sample with already learned samples (table 3). In both cases few-shot learning fails, as the neural network forgets the pre-trained samples, indicating that both ideas are important for the success.

6. Discussion

Two simple ideas lead to successful few-shot learning: disabling learning of the first layer during few-shot learning and adding samples from pre-training to every shot. Both measures could be integrated into a unified learning procedure.

For example by reducing the learning rate of the first layer by a factor of 0.01 the same neural network could be used for pre-training as well as for few-shot learning. Even

shot	accuracy old digits	accuracy new digits	accuracy only new digits	overall accuracy
1	0.569	0.508	0.517	0.558
2	0.672	0.642	0.659	0.661
3	0.565	0.505	0.517	0.557
4	0.649	0.501	0.529	0.626
5	0.564	0.642	0.661	0.584
6	0.684	0.724	0.733	0.696
7	0.364	0.515	0.517	0.398
8	0.458	0.525	0.536	0.474
9	0.454	0.532	0.541	0.477
10	0.451	0.784	0.792	0.508

Table 3: Experimental results as in 1 , but without enriched samples.

shot	accuracy old digits	accuracy new digits	accuracy only new digits	overall accuracy
1	0.896	0.409	0.645	0.809
2	0.754	0.573	0.699	0.722
3	0.926	0.484	0.667	0.841
4	0.855	0.637	0.832	0.816
5	0.836	0.730	0.821	0.813
6	0.890	0.574	0.688	0.830
7	0.935	0.683	0.893	0.880
8	0.879	0.778	0.881	0.860
9	0.878	0.806	0.914	0.865
10	0.874	0.801	0.906	0.862

Table 4: Experimental results with pre-training and few-shot learning using the same neural network with a reduced learning rate of the first layer (by a factor 0.01) .

if it performs a little worse than our original procedure, the few-shot learning succeeds (table 4).

Enriching the few-shot samples with “old samples” can be done straight forward. By keeping one or more samples for each label and use them as “old samples”, it is even possible to use the few-shot learning procedure for pre-training. To test this approach we use the few-shot procedure with shots of one digit together with the last four digits as pre-training followed the same few-shot learning as before (table 5). This way the same procedure is used to learn 8 digits as pre-training, followed by learning the remaining two digits as few-shot learning. Using 5000 shots for pre-training, reaching about 85% accuracy, and 10 shots for few-shot learning we reach more than 75% over all accuracy. Using 100 shots for few-shot learning increases the accuracy to 83%. Further tests indicate, that increasing the number of shots for pre-training increases the few-shot accuracy.

shot	accuracy old digits	accuracy new digits	accuracy only new digits	overall accuracy
1	0.786	0.411	0.827	0.700
2	0.719	0.519	0.753	0.676
3	0.793	0.589	0.880	0.745
4	0.755	0.630	0.873	0.723
5	0.761	0.657	0.882	0.732
6	0.771	0.651	0.884	0.738
7	0.778	0.645	0.883	0.744
8	0.772	0.674	0.852	0.739
9	0.765	0.734	0.888	0.751
10	0.792	0.714	0.888	0.768
100	0.840	0.856	0.963	0.830

Table 5: The few-shot learning results are shown, after the pre-training is done with the few-shot procedure using 5000 shots (mini-batches consisting of one new sample together with the last four used samples), reaching an accuracy of about 85%. This way pre-training and few-shot learning is done the same way. The neural network is the same one used in previous experiments, but the learning rate of the first two layers is decreased by a factor of 0.5 and 0.25 respectively. The first layer is random initialized with a standard derivative of 0.1, the other layers with 0.5.

7. Outlook

The experience from the relatively small problem of hand written digits indicate, that our few-shot procedure is successful. Some tests increasing the size of the hidden layers and the amount of pre-training shows even higher accuracy. Overfitting seems not to be a problem, at least not for hidden sizes up to 1024. Therefore we expect increasing the size of the neural network and increasing the number of labels might increase the accuracy, which would support the use in real world application.

- [1] David Silver, Aja Huang, Christopher J. Maddison, Arthur Guez, Laurent Sifre, George van den Driessche, Julian Schrittwieser, Ioannis Antonoglou, Veda Panneershelvam, Marc Lanctot, Sander Dieleman, Dominik Grewe, John Nham, Nal Kalchbrenner, Ilya Sutskever, Timothy Lillicrap, Madeleine Leach, Koray Kavukcuoglu, Thore Graepel, and Demis Hassabis. Mastering the game of go with deep neural networks and tree search. *Nature*, 529:484–503, 2016.
- [2] David Silver, Julian Schrittwieser, Karen Simonyan, Ioannis Antonoglou, Aja Huang, Arthur Guez, Thomas Hubert, Lucas Baker, Matthew Lai, Adrian Bolton, Yutian Chen, Timothy Lillicrap, Fan Hui, Laurent Sifre, George van den Driessche, Thore Graepel, and Demis Hassabis. Mastering the game of go without human knowledge. *Nature*, 550:354–, October 2017.
- [3] Yaqing Wang, Quanming Yao, James T. Kwok, and Lionel M. Ni. Generalizing from a few examples: A survey on few-shot learning. *ACM Comput. Surv.*, 53(3), June 2020.
- [4] Pierre Baldi and Peter Sadowski. A theory of local learning, the learning channel, and the optimality of backpropagation. *Neural Networks*, 83:51 – 74, 2016.
- [5] Gregory Cohen, Saeed Afshar, Jonathan Tapson, and André van Schaik. Emnist: an extension of mnist to handwritten letters. *arXiv preprint arXiv:1702.05373*, 2017.
- [6] Thomas Kluyver, Benjamin Ragan-Kelley, Fernando Pérez, Brian Granger, Matthias Bussonnier, Jonathan Frederic, Kyle Kelley, Jessica Hamrick, Jason Grout, Sylvain Corlay, Paul Ivanov, Damián Avila, Safia Abdalla, and Carol Willing. Jupyter

notebooks – a publishing format for reproducible computational workflows. In F. Loizides and B. Schmidt, editors, *Positioning and Power in Academic Publishing: Players, Agents and Agendas*, pages 87 – 90. IOS Press, 2016.

- [7] Detlef Schmicker. Towards few-shot learning. https://github.com/dsmic/towards_few_shot_learning, 2020.